

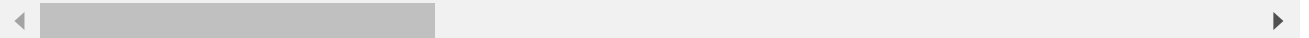
```
import os
import nltk
import nltk.corpus
```

```
nltk.download('punkt')
```

```
[nltk_data] Downloading package punkt to /root/nltk_data...
[nltk_data]   Unzipping tokenizers/punkt.zip.
True
```

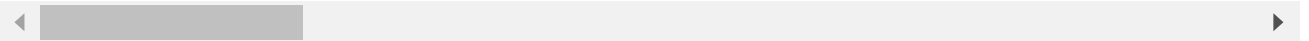
```
from nltk.tokenize import sent_tokenize
text="""Earth Day is an annual event on April 22 to demonstrate support for environmental
tokenized_text=sent_tokenize(text)
print(tokenized_text)
```

```
['Earth Day is an annual event on April 22 to demonstrate support for environmental
```



```
from nltk.tokenize import word_tokenize
tokenized_word=word_tokenize(text)
print(tokenized_word)
```

```
['Earth', 'Day', 'is', 'an', 'annual', 'event', 'on', 'April', '22', 'to', 'demonstr
```




```
from nltk.probability import FreqDist
fdist = FreqDist(tokenized_word)
print(fdist)
```

```
<FreqDist with 41 samples and 45 outcomes>
```

```
fdist.most_common(2)
```

```
[('on', 2), ('April', 2)]
```

```
# Frequency Distribution Plot
import matplotlib.pyplot as plt
fdist.plot(30,cumulative=False)
plt.show()
```



```
nltk.download('stopwords')
```

```
[nltk_data] Downloading package stopwords to /root/nltk_data...
[nltk_data] Unzipping corpora/stopwords.zip.
True
```


```
4      E-      in  ev  str  ipf  rel  ec  f  f  i  -  lu  v  rel  ave  na  ob
```

```
from nltk.corpus import stopwords
stop_words=set(stopwords.words("english"))
print(stop_words)
```

```
{'now', 'he', 've', 'but', 'whom', 'its', 'yours', 'below', 'when', 'all', 'too', 'y
```

```
filtered_sent=[]
for w in tokenized_word:
    if w not in stop_words:
        filtered_sent.append(w)
print("Tokenized Sentence:",tokenized_word)
print("Filterd Sentence:",filtered_sent)
```

```
Tokenized Sentence: ['Earth', 'Day', 'is', 'an', 'annual', 'event', 'on', 'April', '
Filterd Sentence: ['Earth', 'Day', 'annual', 'event', 'April', '22', 'demonstrate',
```




```
# Stemming
from nltk.stem import PorterStemmer
from nltk.tokenize import sent_tokenize, word_tokenize
```

```
ps = PorterStemmer()
```

```
stemmed_words=[]
for w in filtered_sent:
    stemmed_words.append(ps.stem(w))
```

```
print("Filtered Sentence:",filtered_sent)
print("Stemmed Sentence:",stemmed_words)
```

```
Filtered Sentence: ['Earth', 'Day', 'annual', 'event', 'April', '22', 'demonstrate',
Stemmed Sentence: ['earth', 'day', 'annual', 'event', 'april', '22', 'demonstr', 'su
```



```
nltk.download('wordnet')
```

```
[nltk_data] Downloading package wordnet to /root/nltk_data...
```

```
[nltk_data]  Unzipping corpora/wordnet.zip.
True
```

```
#Lexicon Normalization
#performing stemming and Lemmatization
```

```
from nltk.stem.wordnet import WordNetLemmatizer
lem = WordNetLemmatizer()
```

```
from nltk.stem.porter import PorterStemmer
stem = PorterStemmer()
```

```
word = "flying"
print("Lemmatized Word:",lem.lemmatize(word,"v"))
print("Stemmed Word:",stem.stem(word))
```

```
Lemmatized Word: fly
Stemmed Word: fli
```

```
sent = "Hey hii how is your day?"
tokens=nltk.word_tokenize(sent)
print(tokens)
```

```
['Hey', 'hii', 'how', 'is', 'your', 'day', '?']
```

```
nltk.download('averaged_perceptron_tagger')
```

```
[nltk_data] Downloading package averaged_perceptron_tagger to
[nltk_data] /root/nltk_data...
[nltk_data] Unzipping taggers/averaged_perceptron_tagger.zip.
True
```

```
nltk.pos_tag(tokens)
```

```
[('Hey', 'NNP'),
 ('hii', 'VBD'),
 ('how', 'WRB'),
 ('is', 'VBZ'),
 ('your', 'PRP$'),
 ('day', 'NN'),
 ('?', '.')]

```

```
import pandas as pd
from google.colab import files
uploaded = files.upload()
```

train.tsv

- **train.tsv**(n/a) - 8481022 bytes, last modified: 4/22/2021 - 100% done  
Saving train.tsv to train.tsv

```
data=pd.read_csv('train.tsv', sep='\t')
data.head()
```

	PhraseId	SentenceId	Phrase	Sentiment
0	1	1	A series of escapades demonstrating the adage ...	1
1	2	1	A series of escapades demonstrating the adage ...	2
2	3	1	A series	2
3	4	1	A	2
4	5	1	series	2


```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 156060 entries, 0 to 156059
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0   PhraseId    156060 non-null int64
1   SentenceId  156060 non-null int64
2   Phrase      156060 non-null object
3   Sentiment   156060 non-null int64
dtypes: int64(3), object(1)
memory usage: 4.8+ MB
```

```
data.Sentiment.value_counts()
```

```
2    79582
3    32927
1    27273
4     9206
0     7072
Name: Sentiment, dtype: int64
```

```
Sentiment_count=data.groupby('Sentiment').count()
plt.bar(Sentiment_count.index.values, Sentiment_count['Phrase'])
plt.xlabel('Review Sentiments')
plt.ylabel('Number of Review')
plt.show()
```



```
from sklearn.feature_extraction.text import CountVectorizer
from nltk.tokenize import RegexpTokenizer
#tokenizer to remove unwanted elements from out data like symbols and numbers
token = RegexpTokenizer(r'[a-zA-Z0-9]+')
cv = CountVectorizer(lowercase=True,stop_words='english',ngram_range = (1,1),tokenizer = t
text_counts= cv.fit_transform(data['Phrase'])
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(
    text_counts, data['Sentiment'], test_size=0.3, random_state=1)
from sklearn.naive_bayes import MultinomialNB
#Import scikit-learn metrics module for accuracy calculation
from sklearn import metrics
# Model Generation Using Multinomial Naive Bayes
clf = MultinomialNB().fit(X_train, y_train)
predicted= clf.predict(X_test)
print("MultinomialNB Accuracy:",metrics.accuracy_score(y_test, predicted))
```

MultinomialNB Accuracy: 0.6049169122986885

## New Section

✓ 1s completed at 12:36 PM

